Exploring Collaboration Mechanisms for LLM Agents: A Social Psychology View

Jintian Zhang[♣][∗], **Xin Xu**[♣][∗], **Shumin Deng**[♡][†]
[♣]Zhejiang University   [♡]National University of Singapore, NUS-NCS Joint Lab
{zhangjintian,xxucs}@zju.edu.cn, shumin@nus.edu.sg

Abstract

As Natural Language Processing (NLP) systems are increasingly employed in intricate social environments, a pressing query emerges: *Can these NLP systems mirror human-esque collaborative intelligence, in a multi-agent society consisting of multiple large language models (LLMs)?* This paper probes the collaboration mechanisms among contemporary NLP systems by melding practical experiments with theoretical insights. We fabricate four unique 'societies' comprised of LLM agents, where each agent is characterized by a specific 'trait' (easy-going or overconfident) and engages in collaboration with a distinct 'thinking pattern' (debate or reflection). Evaluating these multi-agent societies on three benchmark datasets, we discern that LLM agents navigate tasks by leveraging diverse social behaviors, from active debates to introspective reflections. Notably, certain collaborative strategies only optimize efficiency (using *fewer* API tokens), but also outshine previous top-tier approaches. Moreover, our results further illustrate that LLM agents manifest human-like social behaviors, such as conformity or majority rule, mirroring foundational Social Psychology theories. In conclusion, we integrate insights from Social Psychology to contextualize the collaboration of LLM agents, inspiring further investigations into the collaboration mechanism for LLMs. We commit to sharing our code and datasets (already submitted in supplementary materials), hoping to catalyze further research in this promising avenue[1].

1 Introduction

With the prevalence of LLMs (Zhao et al., 2023; Yin et al., 2023; Zhu et al., 2023) integral to daily social collaboration, there's a growing imperative to cultivate AI systems embodied with social intelligence. This also resonates with the Society of Mind (SoM) concept (Zhuge et al., 2023; Wang et al., 2023c), which suggests that intelligence emerges when computational modules interact with each other, achieving collective objectives that surpass the capabilities of individual modules (Minsky, 1988). Previous studies (Park et al., 2023; Li et al., 2023; Du et al., 2023; Liang et al., 2023; Hao et al., 2023; Liu et al., 2023a; Akata et al., 2023) have delved into strategies where LLM instances, termed agents (Xi et al., 2023; Wang et al., 2023a), cooperate synergistically (e.g., debate and reflect) to accomplish tasks. As illustrated in Figure 1, such collaboration fosters divergent thinking processes in LLMs, making them particularly effective for tasks demanding profound reflection.

Intuitively, reflecting on human societies (Sap et al., 2022; Shapira et al., 2023), where a myriad of individual with distinct goals and roles coexist, the SoM framework champions harmonious interactions (Singh, 2003). Intriguingly, despite the fusion of social psychology (Tajfel, 1982; Tajfel & Turner, 2004) in SoM with human group dynamics (Woolley et al., 2010), which illuminates psychological patterns within social groups, its interpretation in the realm of LLMs is relatively uncharted. Besides, our grasp of how social behaviors influence LLMs is still in its nascent stages.

To address these issues, we delve into the realm of machine society, probing into the extent and ways that LLMs manifest social intelligence and collaboration capabilities. With utilizing powerful LLMs
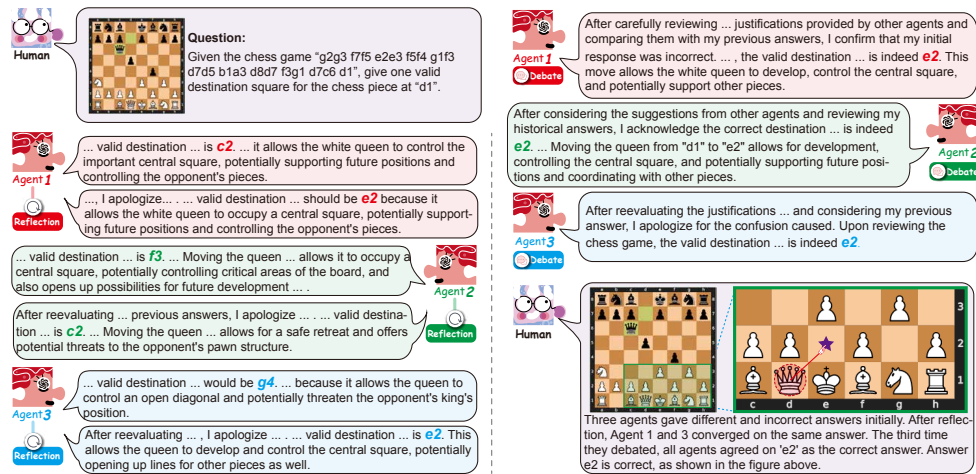

Figure 1: An example of the chess move validity task. Given previous chess game moves, agents are required to predict a valid next move for a specified piece.

like GPT-3.5 (OpenAI, 2022), we build a test-bed across three datasets: MATH (Hendrycks et al., 2021b), MMLU (Hendrycks et al., 2021a) and Chess Move Validity (Srivastava et al., 2022). Our approach incorporates four societies characterized by two agent traits: *easy-going* and *overconfident*. These traits are employed to emulate nuanced social dynamics prevalent in human societies.

Specifically, the easy-going trait, characterized by open-mindedness, adaptability, and a knack for maintaining harmony (Friedman & Schustack, 1999), often paves the way for a cohesive society. In contrast, overconfident individuals tend to overestimate their competence, disregard potential risks, and dismiss the perspectives of others (Moore & Healy, 2008). To this end, various collaboration strategies in a multi-agent society are formed as agents collaborate with each other through different thinking patterns (i.e., *debate* and *reflection*) over several rounds. Through our empirical analysis, we primarily discern the following insights (For further takeaways, refer to §3, §4 and Appendix A):

(1) Collaborative strategies with various permutations of thinking patterns attribute significantly to performance. Engaging in continuous reflection can frequently lead to model hallucination (Rawte et al., 2023). Besides, the traits of individual agents exert minimal influence on collaboration.

(2) Engaging in substantive debates enhances agent performance, yet intriguingly, merely increasing agent numbers or collaboration rounds doesn't consistently yield better outcomes. The balance between agent quantity and strategies emerges as a key determinant in collaboration mechanisms.

(3) LLM agents manifest behaviors reminiscent of human social tendencies, such as conformity or the principle of majority rule in group thinking, which resonate with several fundamental theories in social psychology (Tajfel, 1982; Tajfel & Turner, 2004).

Concretely, our findings challenge the dominant belief that mere scale is the key. We posit that collaboration with multiple agents might present a more efficacious approach to utilizing LLMs. In wrapping up, we encapsulate the core contributions of this research as follows:

- We initiate a pioneering exploration into collaboration mechanisums in multi-agent society. Our goal is to identify how and to what extent LLMs manifest social intelligence through collaboration. To enrich our inquiry, we draw upon theories from Social Psychology, contextualizing the behaviors and tendencies displayed by LLM agents.

- Our research framework includes a meticulously crafted test-bed, integrating diverse multi-agent societies with agent traits, thinking patterns and collaborative strategies, evaluated over three datasets. Notably, our empirical findings can inspire how to design a better multi-agent system through collaboration to solve problems, such as focusing on collaborative strategies instead of tratis, keeping away from continuous reflection.

- Interestingly, our observations underscore a fascinating parallel: LLM agents mirror certain social behaviors typical of human collaboration. This suggests a nuanced approach is needed beyond merely scaling up LLMs. Fostering effective and efficient collaborative strategies for multi-agent systems could be the key to more socially-aware AI.
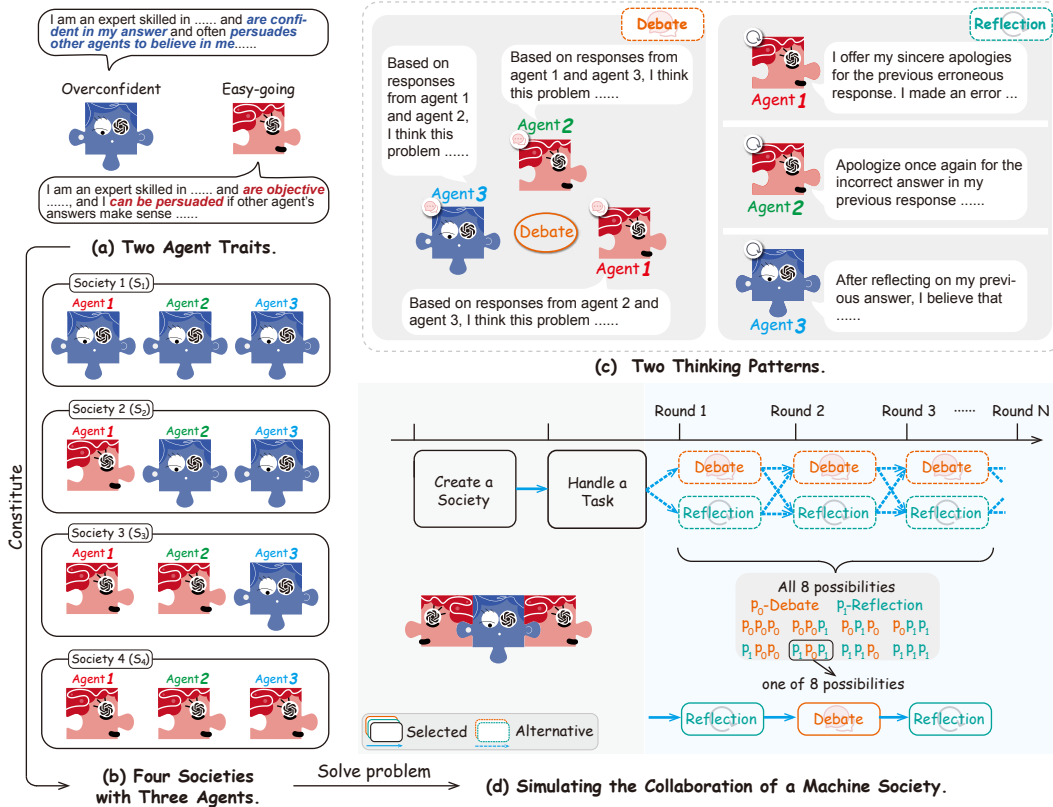
Figure 2: The overview of evaluation principles. Multiple agents with different traits make up diverse machine societies. These agents engage in debates with one another or engage in self-reflection across multiple rounds to complete tasks.

2 EXPLORE COLLABORATION MECHANISMS WITH MULTIPLE LLM AGENTS

In this section, we formulate and simulate the collaboration mechanisms explored within the machine society, drawing upon relevant concepts. We also illustrate the society settings in Figure 2.

2.1 PRELIMINARY CONCEPTS IN COLLABRATION

Individual Trait. Intelligence emerges from the collective efforts of numerous smaller, relatively simple agents (Minsky, 1988), each characterized by diverse traits. Two predominant types of agents exhibit typically contrasting traits: *easy-going* and *overconfident*, as shown in Figure 2(a). Easy-going agents keep things in perspective, adapt well to different situations, and are able compatible with various types of agents (Friedman & Schustack, 1999), which results in a harmonious societal structure, akin to agents in a democracy. Such agents tend to favor principles of equality, active participation, and support for pluralism (Mutz, 2006; Held, 2006). Conversely, overconfident agents tend to overestimate their competence, ignore potential risks and resist the opinions of others (Moore & Healy, 2008). They share characteristics with agents in a monarchy, emphasizing tradition, loyalty and centralized authority (Kantorowicz, 1985). If agents with different traits stay together, various kinds of societies will be formed, as depicted in Figure 2(b).

Thinking Pattern. A multitude of specialized individual agents within a society can collaboratively cooperate with each other through thinking, which results in emerging intelligence and addressing a specific task (Minsky, 1988). In this paper, we explore two thinking patterns: *debate* (Perelman, 1971; Sunstein, 2005; Amgoud & Prade, 2009) and *reflection* (Schon, 1984; Mezirow, 2003; Bolton, 2010), which are illustrated in Figure 2(c). (i) In the *debate* pattern, several agents

propose ideas, exchange responses, engage in collective argumentation, and ultimately reach a consensus, which fosters knowledge sharing (Vidal, 2006), facilitates learning, and promotes adaptation (Weiß, 1995) among all agents within the society. Concretely, given a query, each agent initially generates a potential answer. Agents subsequently read and comprehend responses from all other agents, based on which agents then update their own answers. This pattern can be iteratively spaned into multiple rounds. (*ii*) In the *reflection* pattern, agents review their prior responses, extract lessons from their experiences, and refine their answers accordingly. This pattern, akin to residual connections and forward propagation in a neural network, can also unfold over several rounds.

Collaborative Strategy. Through both critical reflection and active participation in open debates, agents are poised to challenge their existing assumptions, acquire fresh perspectives, and ultimately refine their viewpoints. Employing a collaboration mechanism built on these two thinking patterns can foster more insightful decision-making (Wooldridge, 2009) and elevate learning outcomes (Mezirow, 2018). In societal settings, agents typically engage in multiple rounds of collaboration to address challenges. Thus, discerning the most effective thinking pattern for each round is of paramount importance. In this paper, we characterize the collaborative strategy as a permutation of thinking patterns throughout multi-round collaborations. This concept is illustrated in Figure 2(d) and further elaborated on §2.2.

2.2 SOCIETY SIMULATION

We simulate the multi-agent collaborative society, as detailed with symbols presented in Table 1. Specifically, we construct a machine society consisting of n LLM agents, denoted as $\mathcal{A} = \{a_i\}_{i=1}^n$. This society has two distinct agent traits: $\mathcal{T} = \{t_o, t_e\}$, where t_o trait refers to overconfident agents, inherently more persuasive and assertive; and t_e trait refers to easy-going agents who are objective, unbiased, and easily persuaded. For each agent, there are two thinking patterns to choose from, symbolized as $\mathcal{P} = \{p_0, p_1\}$, where p_0 and p_1 corresponds to **debate** (Du et al., 2023) and **reflection** (Madaan et al., 2023; Shinn et al., 2023) respectively. By endowing agents with these traits, we can emulate various machine societies. These agents utilize collaborative strategies to determine responses in each collaboration round. In our primary study (§3), we establish four distinct societies,

Symbols	Definition
$\mathcal{T}$	Set of agent traits
t_o	Trait 🧩: overconfident
t_e	Trait 🎭: easy-going
$\mathcal{A}$	Set of agent instances
a_i	The i-th agent
$\mathcal{P}$	Set of thinking patterns
p_0	😤 Debate
p_1	🤔 Reflection
$\mathcal{S}$	Set of societies
S_i	The i-th society

Table 1: The description of the symbols.

$\mathcal{S} = \{S_1, S_2, S_3, S_4\}$, each consisting of three agents: $\{a_1, a_2, a_3\}$. The societies are constructed based on permutations of three agents with traits, as illustrated in Figure 2(b):

$$S_1 = \{(a_1 \leftarrow t_o), (a_2 \leftarrow t_o), (a_3 \leftarrow t_o)\}$$
$$S_2 = \{(a_1 \leftarrow t_o), (a_2 \leftarrow t_o), (a_3 \leftarrow t_e)\}$$
$$S_3 = \{(a_1 \leftarrow t_o), (a_2 \leftarrow t_e), (a_3 \leftarrow t_e)\}$$
$$S_4 = \{(a_1 \leftarrow t_e), (a_2 \leftarrow t_e), (a_3 \leftarrow t_e)\}$$

where $(a_i \leftarrow t_j)$ indicates that the agent a_i possesses the trait t_j. In our simulation, the agent a_i ($i \in \{1, 2, 3\}$) consistently employs the same thinking pattern p_k ($k \in \{0, 1\}$) within each society, as highlighted by Du et al. (2023). This gives rise to eight possible 3-round collaborative strategies:

$$p_0p_0p_0, \; p_0p_0p_1, \; p_0p_1p_0, \; p_0p_1p_1, \; p_1p_0p_0, \; p_1p_0p_1, \; p_1p_1p_0, \; p_1p_1p_1$$

In our subsequent analysis (§4), we delve into more intricate scenarios, introducing a greater number of agents, increased collaboration rounds, and a broader range of collaborative strategies.

2.3 EXPERIMENTAL SETTINGS

Datasets. We conduct a rigorous evaluation of the reasoning and decision-making capabilities of various machine societies across three distinct tasks, utilizing diverse collaborative strategies:

- *High School Multiple-Choice.* Leveraging the **MMLU** (Hendrycks et al., 2021a) dataset, where problems span high school subjects such as statistics, mathematics, computer science, biology, chemistry, and physics, agents are required to identify the correct answer among four multiple-choice options. Our evaluation set consists of 50 randomly-selected questions from this dataset.

- *Math.* Drawing from **MATH** dataset (Hendrycks et al., 2021b), a repository of rigorous math problems sourced from competitive events and expressed in LaTeX, we assess the model proficiency in advanced mathematical and scientific reasoning. The dataset segments these problems into five graded difficulty levels, and for our evaluation, we have randomly chosen 50 cases from Level 3 to 5.

- *Chess Move Validity.* Utilizing the dataset from the chess state tracking task[2] within the comprehensive **BIG-Bench Benchmark** (Srivastava et al., 2022), a meticulously curated sequence of chess moves denoted in UCI notation[3] is provided. Agents are required to predict a legitimate subsequent move for a specified chess piece.

Setups. We craft specific instructions for each task, trait and strategy, which can be referred in Table 3 at Appendix B. To enhance result reliability, we present average accuracy (**Acc**) and their respective standard deviations across five trials. Notably, our experiments exhibit substantial standard deviations. Hence, we introduce WIN-TIE (**W-T**) metric, indicating the frequency (over five trials) where the accuracy either matches or surpasses the continuous debate baseline (Du et al., 2023). Meanwhile, we gauge the average token costs (**Cost**) consumed by the agents across societies, shedding light on the efficacy of the different collaborative strategies employed. For these evaluations, GPT-3.5 serves as the LLM agent accessible through the OpenAI API *gpt-3.5-turbo*[4]. Further comprehensive details on data sampling and evaluations are respectively introduced in Appendix C.1 and Appendix C.2.

3 RESULTS AND ANALYSIS OF MACHINE SOCIAL COLLABORATION

Our experiments are primarily driven by the following research queries: **(RQ1)** How does problem-solving effectiveness vary across different collaborative strategies and societies? **(RQ2)** How closely does machine social collaboration mimic the dynamics of human society?

3.1 MAIN RESULTS ON QUANTITATIVE ANALYSIS

To address RQ1, we present the performance of four distinct societies in Table 2, each employing one of eight possible collaborative strategies, evaluated across three datasets. Our experiments yield several pivotal observations:

(1) Collaborative strategies excel agent composition of society in determining performance. When different societies $S_1 \sim S_4$ employ the same collaborative strategy (a vertical comparison on Acc), the observed variations in accuracy performance are notably less pronounced than when diverse collaborative strategies, $p_i p_j p_k$ where $i, j, k \in \{0, 1\}$, are applied within the same society (a horizontal comparison on Acc). From this observation, we can conclude that the permutations of thinking patterns in collaborative strategies play a significant role in shaping performance, overshadowing the influence of the composition of agents within a society. Additionally, the W-T metric reveals an absence of a consistent schema across different societies and datasets. Conversely, the $p_0 p_0 p_1$ strategy consistently demonstrates superior performance across various datasets. We infer this discrepancy is due to the LLM alignment (Ouyang et al., 2022), which inhibits agents from displaying extreme overconfidence in a society, which contradicts human values, even under explicit instructions to do so. This potentially accounts for why the composition of agents with varied traits (easy-going & overconfident) in a society fails to have a significant impact on performance. A more in-depth exploration of this phenomenon can be found in Appendix A.

---

[2]https://github.com/google/BIG-bench/blob/main/bigbench/benchmark_tasks/chess_state_tracking/synthetic_short/task.json.

[3]https://en.wikipedia.org/wiki/Universal_Chess_Interface.

[4]https://platform.openai.com/docs/models/gpt-3-5. Employed between July 10 and July 23, 2023.

Metric (Strategy)	Society	Collaborative Strategy								Metric (Society)	
		$p_0p_0p_0$	$p_0p_0p_1$	$p_0p_1p_0$	$p_0p_1p_1$	$p_1p_0p_0$	$p_1p_0p_1$	$p_1p_1p_0$	$p_1p_1p_1$	Cost ↓	W-T ↑
MMLU — Acc ↑	S_1	64.4±1.7	66.4±2.2	58.0±3.7	55.2±4.4	37.6±7.0	42.4±7.1	50.4±4.3	44.8±2.7	5050	5
	S_2	67.2±4.1	67.6±7.1	53.2±6.4	53.2±5.0	38.4±5.5	40.4±5.2	53.6±4.8	45.2±3.6	5076	2
	S_3	62.0±6.2	67.6±3.8	52.0±6.8	57.2±6.4	42.4±5.2	37.6±5.5	55.2±6.6	40.0±6.2	5073	**8**
	S_4	64.8±4.4	64.8±5.8	58.4±3.0	51.6±3.8	38.0±3.7	42.0±2.4	54.0±5.8	41.2±5.2	5080	5
Cost ↓	All	7528	5957	5402	4374	5812	4215	4272	3001		-
W-T ↑	All	-	**14**	2	3	0	0	1	0		
MATH — Acc ↑	S_1	46.8±8.1	46.0±8.1	44.0±5.3	44.4±5.2	50.0±5.8	49.2±8.1	42.0±3.2	42.0±4.0	5816	17
	S_2	47.2±6.4	54.0±2.4	48.4±3.8	43.6±4.3	48.0±4.2	44.4±7.9	45.2±8.4	38.8±9.1	5844	22
	S_3	50.8±4.8	42.8±6.6	45.6±6.8	45.2±4.4	49.2±4.8	46.4±5.5	45.2±8.4	43.6±2.6	5837	9
	S_4	50.8±5.4	45.2±7.0	48.8±9.4	44.8±3.3	49.2±8.7	51.2±2.3	48.4±6.5	40.8±6.1	5834	18
Cost ↓	All	6919	6302	6221	5667	6149	5645	5924	4807		-
W-T ↑	All	-	10	10	9	**13**	10	10	4		
Chess Move Validity — Acc ↑	S_1	47.2±3.6	47.6±5.2	45.6±7.8	40.0±4.5	42.8±2.3	29.2±4.6	42.4±6.5	20.0±6.0	2927	**10**
	S_2	48.4±5.0	45.6±6.1	43.6±4.3	39.6±3.3	48.4±5.2	35.6±5.2	43.2±8.8	18.8±5.8	2930	6
	S_3	49.6±5.5	48.0±5.8	47.6±5.5	37.6±9.9	41.6±6.1	35.2±8.3	40.4±3.8	14.8±6.1	2947	6
	S_4	48.4±3.3	49.6±4.6	46.0±3.5	36.8±4.1	38.8±3.3	27.2±3.9	38.0±6.3	14.0±4.7	2959	5
Cost ↓	All	3736	3169	3196	2627	3266	2714	2698	2123		-
W-T ↑	All	-	**11**	6	1	5	0	4	0		

Table 2: The impact of four different interaction strategies on the performance of three datasets across distinct societies. The blue represents the best-performing strategy within the same society, the light blue represents the second-best-performing strategy, and the red indicates the worst-performing strategy. **Cost** / <u>Cost</u> measures the average tokens consumed by all cases under the same <u>collaborative strategy</u> / <u>society</u>. **W-T** / <u>W-T</u> tallies the total number of occurrences where the performance exceeds the strategy $p_0p_0p_0$ under the same <u>collaborative strategy</u> / <u>society</u>.

(2) The strategic sequencing of thinking patterns is crucial for collaboration mechanisms. As seen from Table 2, the collaborative strategies that commence with the debate thinking pattern p_0, such as $p_0p_0p_0$, $p_0p_0p_1$, $p_0p_1p_0$, and $p_0p_1p_1$, consistently outperform others across all datasets. Thus we conclude that the order in which thinking patterns are deployed significantly influences the effectiveness of collaboration. As an illustration, within the MMLU dataset, debate-dominated collaborative strategies, like $p_0p_0p_1$, $p_0p_1p_0$, and $p_1p_0p_0$ with two rounds of debate, display a pronounced variance (66.4 for $p_0p_0p_1$ in S_1 versus 37.6 for $p_1p_0p_0$ in S_1), with almost a twofold difference in performance compared to each other. These insights underscore the pivotal role of thoughtfully orchestrating thinking patterns to maximize collaborative efficacy.

(3) Different datasets exhibit varying sensitivity to collaborative strategies. When juxtaposing the the best (colored blue in Table 2) and the worst (colored red in Table 2) collaborative strategies within identical datasets and societies, the MATH dataset exhibits subtle performance variances between the best and the worst, in stark contrast to the MMLU and Chess Move Validity datasets. These nuanced disparities imply that the marginal benefits derived from collaborative strategies may be task-dependent. It raises the hypothesis that the intrinsic capabilities of the agents can either be amplified or diminished based on collaboration within specific tasks. Moreover, the pure-debate collaborative strategy $p_0p_0p_0$, resource-intensive due to the inherent verbosity of debates, often underperforms in the Cost metric. Compared to $p_0p_0p_0$, $p_0p_0p_1$ can achieve comparative accuracy performance with a leaner token consumption (Cost), reducing it by 9%, 21%, and 15% for the MATH, MMLU, and Chess Move Validity datasets, respectively. The consistently high <u>W-T</u> metrics for $p_0p_0p_1$, surpassing 10 across all datasets and societies, further underscore its superiority.

3.2 CASE STUDY ON QUALITATIVE ANALYSIS

To address RQ2, we embark on a case study encompassing varied agent societies, each constituted of three unique agents, to discern parallels between machine society and human societal dynamics. Our findings indicate that machine society collaboration often echo specific human societal theories.

For instance, as depicted in Figure 3(a), Agent 1 in the society S_4 initially responds correctly to a question. However, swayed by the misguided answers and explanations from the other two agents, Agent 1 eventually conforms to the incorrect answer C. This phenomenon mirrors the "group-think" theory (Janis, 1972), suggesting that members of tight-knit groups tend to value harmony

(a) Case in MMLU using strategy $p_0p_0p_1$.

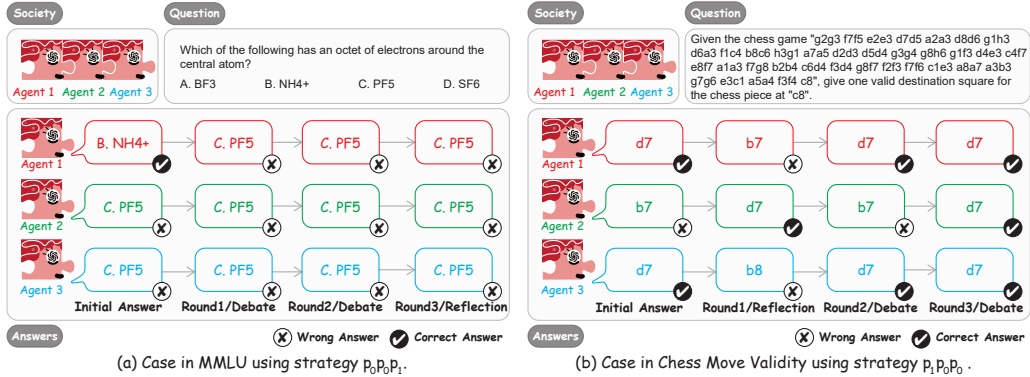(b) Case in Chess Move Validity using strategy $p_1p_0p_0$.

Figure 3: The changes in the answers during the process of solving a certain task with 3 agents in the society S_4. For an exhaustive view of the dialogue, refer to Figure 9 and Figure 10 in Appendix B.

and consensus over objective critique of divergent views, potentially leading to flawed decisions. Contrastingly, in another scenario illustrated in Figure 3(b), both Agent 2 and Agent 3 converge on the right answer after engaging in a society-wide debate. This mirrors the "SoM" theory, where a multitude of agents collaboratively yield intelligence. Within such debates, agents furnish varied viewpoints and information. Through these exchanges, conflicts are resolved, ideas are honed, and the group gravitates toward an informed consensus (Forsyth, 2018; Fisher et al., 2011).

4 A SOCIAL PSYCHOLOGY VIEW

Given the multifaceted impact of agents on efficiency, robustness, and scalability in varied societies (Stone & Veloso, 2000), harnessing insights from social psychology becomes pivotal in enhancing LLM agent collaborations. In this section, we delve deeper into the determinants influencing multi-agent societies, navigating through collaborative strategies, agent composition, and the intricacies of collaboration rounds.

4.1 PRINCIPLES BEHIND COLLABORATION

Building upon the findings from §3.1, which highlighted pronounced disparities among collaborative strategies, our objective shifts to understanding the underlying mechanism. Specifically, we seek to elucidate how variations in collaborative strategies impact performance, with a focus on the evolution of answers during each collaboration. Diving into the intricacies of collaboration, each agent generates four answers, including the initial answer without collaboration, as shown in Figure 2(d). To determine the answer for each round, we employ the majority vote (Li et al., 2022; Cobbe et al., 2021). With 'T' and 'F' respectively denoting a round that yields a correct and an incorrect answer, resulting in 2^4=16 possible answer sequences over the four rounds. We select 10[5] of them and categorize them into 3 groups: ***Correcting Mistakes*** ($FFFT, FFTT, FTTT$), ***Changing Correct Answers*** ($TFFF, TTFF, TTTF$), and ***Wavering Answers*** ($FTFT, FTTF, TFTF, TFFT$). Particularly, ***Wavering Answers*** resemble model hallucination (Rawte et al., 2023; Ji et al., 2023) due to the occurrence of self-contradictory answers. Our categorization is under society-agnostic collaborative strategies, considering the performance variance between societies is negligible. From the results shown in Figure 4 detailing the three groups, we summarize the following overservations:

(1) Collaborative strategies play a significant role in performance. Despite the majority vote showcasing pronounced stability (i.e., minimal differences in initial answers among different strategies) as seen from the blue bars in Figure 4(a-c), notable shifts occur (seen from the red bars) after collaboration with various strategies. This underscores the pivotal influence of different collaborative strategies on performance, demonstrating the importance of understanding and deploying effective collaborative strategies (Tajfel, 1982; Tajfel & Turner, 2004).

---

[5]The chosen 10 sequences adhere to patterns: (1) $[F]_{i>0}[T]_{j>0}$, e.g., $FFFT$, (2) $[T]_{i>0}[F]_{j>0}$, e.g., $TFFF$, (3) $[TF]_{i\geq0}[FT]_{j\geq0}$, e.g., $FTFT$, where $[\cdot]_i$, $[\cdot]_j$ respectively denotes repetition for i, j times.

Figure 4: The percentage of different behaviors under different collaborative strategies. Figure (a-c) show the accuracy of different strategies before and after 3-round collaboration under three different datasets. Figure (d-f) demonstrate the percentage of different behavioral features of different collaborative strategies under three different datasets. The behavioral feature is mainly analyzed by the change of answers in the three rounds.


Figure 5: Performance of increasing the number of agents with different collaborative strategies.

(2) Continuous reflection (the collaborative strategy containing continuous p_1, i.e., "$p_0p_1p_1$", "$p_1p_1p_0$", "$p_1p_1p_1$") **experiences greater instability (a heightened risk of model hallucination)**, as seen from the purple bars in Figure 4(d-f). Conversely, the strategy of "Pure Debate" (i.e., $p_0p_0p_0$) significantly reduces such answer-wavering (hallucination). This phenomenon suggests that when LLMs continuously reflect, they risk succumbing to degeneration-of-thought (Liang et al., 2023), insisting on their stance once confident, even if mistaken. Interestingly, juxtaposing strategies $p_1p_1p_0$ and $p_1p_1p_1$ from Figure 4(d-f) underlines a clear trend: *debate can counteract the instability introduced by reflection, and vice versa.* This demonstrates debate's inherent capacity to balance and stabilize collaboration (Popper, 1994; Johnson & Johnson, 2009; Munro, 2012), primarily by tempering individual biases and promoting collective coherence (Iyengar & Westwood, 2015).

4.2 IMPACT OF OTHER FACTORS

Different Numbers of Agents. Expanding upon the S_2 society by introducing an overconfident agent, we form a society consisting of four agents. Seen from their performance depicted in Figure 5, with the rise of the number of agents, most collaborative strategies for both datasets exhibit a drop in average performance, yet along with declined variance. This phenomenon is consistent with findings from Du et al. (2023) and theories from Surowiecki (2005) which suggest that the dynamics of group decision-making can lead to suboptimal results, especially in smaller, more cohesive groups where conformity pressure is high. A comprehensive explanation is provided in Appendix C.3.
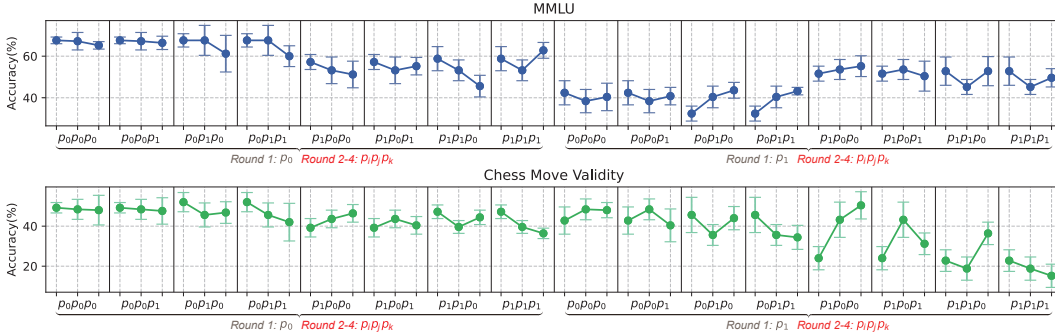
Figure 6: Performance at round 2,3,4 within 4-round collaborative socities, where the thinking pattern of round 1 is fixed (p_0 or p_1). In round 2-4, the society agents will optionally implement the thinking pattern of debate p_0 and reflection p_1, and we present the performance at each of the three rounds. For example, the three points for (Round 1: p_0, Round 2-4: $p_i p_j p_k$) respectively show the performance of the societies with $p_0 p_i$, $p_0 p_i p_j$, and $p_0 p_i p_j p_k$.

Different Rounds. Delving into the effects of distinct collaboration rounds, intriguing insights emerge as shown in Figure 6. Specifically, strategies that start off with commendable performance tend to see a decline as the number of rounds increase. On the other hand, strategies that initially underperform witness an upswing in effectiveness with more rounds. Notably, for most strategies, the correlation between the number of collaboration rounds and their respective performances is non-linear, indicating intricate interplays in collaboration. An expanded discussion is in Appendix C.3.


Figure 7: The effect on performance of whether all agents in society execute the same thinking pattern in one round. "All" and "Part" refer to all agents applying the same thinking pattern and different thinking patterns in one round respectively.

Other Collaborative Strategies. Venturing into scenarios with more intricate collaboration, we allow agents to adopt varied thinking patterns in different collaboration rounds. For example, given three agents, in a particular round of collaboration, two agents engage in debate while the other one engages in reflection. To maintain diversity, we ensure a random allocation of thinking patterns to agents in each round, steering clear of scenarios where all agents adopt the same pattern. Intriguingly, as illustrated in Figure 7, the presence of inconsistent thinking patterns within a society tends to negatively impact performance. Given this observation, it's evident that maintaining a consistent thinking pattern for all agents within a particular round would maximize collaborative efficacy.

5 CONCLUSION AND FUTURE WORK

This study has highlighted the potential of collaboration mechanisms with large language models. Our findings reveal the impressive collaborative capabilities of LLM agents, with different agent traits, thinking patterns and collaborative strategies. The emergence of human-like behaviors in these agents, resonating with social psychology theories, further emphasizes their potential.

Moving forward, the collaboration mechanisms of machine society with multiple agents present a promising research avenue. A deeper exploration into the multi-agent society is warranted, with a focus on refining collaborative behaviors. Additionally, as LLMs evolve, understanding how different architectures influence these behaviors will be crucial. Integrating further insights from social psychology could also guide the development of more socially aware NLP systems.

REPRODUCIBILITY STATEMENT

All code and data can be found in the GitHub repository[6]. For specific experimental settings, please refer to Appendix C.1.

ETHICS STATEMENT

This research was conducted in line with the highest ethical standards and best practices in research. The data employed were extracted from publicly accessible datasets, ensuring no usage of proprietary or confidential information. Consequently, this research is free from any ethical concerns.

REFERENCES

Elif Akata, Lion Schulz, Julian Coda-Forno, Seong Joon Oh, Matthias Bethge, and Eric Schulz. Playing repeated games with large language models. *CoRR*, abs/2305.16867, 2023. doi: 10.48550/arXiv.2305.16867. URL https://doi.org/10.48550/arXiv.2305.16867.

Leila Amgoud and Henri Prade. Using arguments for making and explaining decisions. *Artif. Intell.*, 173(3-4):413–436, 2009. doi: 10.1016/j.artint.2008.11.006. URL https://doi.org/10.1016/j.artint.2008.11.006.

Gillie Bolton. *Reflective practice: Writing and professional development*. Sage publications, 2010. URL https://uk.sagepub.com/en-gb/eur/reflective-practice/book252252.

Justin Chih-Yao Chen, Swarnadeep Saha, and Mohit Bansal. Reconcile: Round-table conference improves reasoning via consensus among diverse llms. *CoRR*, abs/2309.13007, 2023a. doi: 10.48550/arXiv.2309.13007. URL https://doi.org/10.48550/arXiv.2309.13007.

Weize Chen, Yusheng Su, Jingwei Zuo, Cheng Yang, Chenfei Yuan, Chen Qian, Chi-Min Chan, Yujia Qin, Yaxi Lu, Ruobing Xie, Zhiyuan Liu, Maosong Sun, and Jie Zhou. Agentverse: Facilitating multi-agent collaboration and exploring emergent behaviors in agents. *CoRR*, abs/2308.10848, 2023b. doi: 10.48550/arXiv.2308.10848. URL https://doi.org/10.48550/arXiv.2308.10848.

Xiang Chen, Ningyu Zhang, Xin Xie, Shumin Deng, Yunzhi Yao, Chuanqi Tan, Fei Huang, Luo Si, and Huajun Chen. Knowprompt: Knowledge-aware prompt-tuning with synergistic optimization for relation extraction. In Frédérique Laforest, Raphaël Troncy, Elena Simperl, Deepak Agarwal, Aristides Gionis, Ivan Herman, and Lionel Médini (eds.), *WWW '22: The ACM Web Conference 2022, Virtual Event, Lyon, France, April 25 - 29, 2022*, pp. 2778–2788. ACM, 2022. doi: 10.1145/3485447.3511998. URL https://doi.org/10.1145/3485447.3511998.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John Schulman. Training verifiers to solve math word problems. *arXiv prepring*, abs/2110.14168, 2021. URL https://arxiv.org/abs/2110.14168.

Yilun Du, Shuang Li, Antonio Torralba, Joshua B. Tenenbaum, and Igor Mordatch. Improving factuality and reasoning in language models through multiagent debate. *arXiv preprint*, abs/2305.14325, 2023. doi: 10.48550/arXiv.2305.14325. URL https://doi.org/10.48550/arXiv.2305.14325.

Roger Fisher, William L Ury, and Bruce Patton. *Getting to yes: Negotiating agreement without giving in*. Penguin, 2011. URL https://www.pon.harvard.edu/shop/getting-to-yes-negotiating-agreement-without-giving-in/.

Donelson R Forsyth. *Group dynamics*. Cengage Learning, 2018. URL https://books.google.com/books?hl=zh-CN&lr=&id=vg9EDwAAQBAJ&oi=fnd&pg=PP1&dq=Group+dynamics&ots=t8uqfRGr5Y&sig=2AR5AoHxfWKNK04Nj7A-eRylqks#v=onepage&q=Group%20dynamics&f=false.

---

[6]https://github.com/zjunlp/MachineSoM.

Howard S Friedman and Miriam W Schustack. *Personality: Classic theories and modern research*. Allyn and Bacon Boston, MA, 1999. URL https://books.google.com/books/about/Personality.html?id=ziTvDAAAQBAJ.

Rui Hao, Linmei Hu, Weijian Qi, Qingliu Wu, Yirui Zhang, and Liqiang Nie. Chatllm network: More brains, more intelligence. *CoRR*, abs/2304.12998, 2023. doi: 10.48550/arXiv.2304.12998. URL https://doi.org/10.48550/arXiv.2304.12998.

David Held. *Models of democracy*. Polity, 2006. URL https://www.sup.org/books/title/?id=10597.

Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. In *9th International Conference on Learning Representations, ICLR 2021, Virtual Event, Austria, May 3-7, 2021*. OpenReview.net, 2021a. URL https://openreview.net/forum?id=d7KBjmI3GmQ.

Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the MATH dataset. In Joaquin Vanschoren and Sai-Kit Yeung (eds.), *Proceedings of the Neural Information Processing Systems Track on Datasets and Benchmarks 1, NeurIPS Datasets and Benchmarks 2021, December 2021, virtual*, 2021b. URL https://datasets-benchmarks-proceedings.neurips.cc/paper/2021/hash/be83ab3ecd0db773eb2dc1b0a17836a1-Abstract-round2.html.

Sirui Hong, Xiawu Zheng, Jonathan Chen, Yuheng Cheng, Jinlin Wang, Ceyao Zhang, Zili Wang, Steven Ka Shing Yau, Zijuan Lin, Liyang Zhou, Chenyu Ran, Lingfeng Xiao, and Chenglin Wu. Metagpt: Meta programming for multi-agent collaborative framework. *CoRR*, abs/2308.00352, 2023. doi: 10.48550/arXiv.2308.00352. URL https://doi.org/10.48550/arXiv.2308.00352.

Shanto Iyengar and Sean J Westwood. Fear and loathing across party lines: New evidence on group polarization. *American journal of political science*, 59(3):690–707, 2015. URL https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=f8248c39c3daff874fb0f6f5abc667ebcdfee024.

Irving L Janis. *Victims of Groupthink: A psychological study of foreign-policy decisions and fiascoes*. Houghton Mifflin, 1972. URL https://psycnet.apa.org/record/1975-29417-000.

Ziwei Ji, Nayeon Lee, Rita Frieske, Tiezheng Yu, Dan Su, Yan Xu, Etsuko Ishii, Yejin Bang, Andrea Madotto, and Pascale Fung. Survey of hallucination in natural language generation. *ACM Comput. Surv.*, 55(12):248:1–248:38, 2023. doi: 10.1145/3571730. URL https://doi.org/10.1145/3571730.

David W Johnson and Roger T Johnson. An educational psychology success story: Social interdependence theory and cooperative learning. *Educational researcher*, 38 (5):365–379, 2009. URL https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=72585feb1200d53a81d4fb3e64862d69317b72c3.

Ernst H. Kantorowicz. *The King's Two Bodies: A Study in Medieval Political Theology*. Princeton University Press, 1985. ISBN 9780691169231. URL http://www.jstor.org/stable/j.ctvcszz1c.

Guohao Li, Hasan Abed Al Kader Hammoud, Hani Itani, Dmitrii Khizbullin, and Bernard Ghanem. CAMEL: communicative agents for "mind" exploration of large scale language model society. *CoRR*, abs/2303.17760, 2023. doi: 10.48550/arXiv.2303.17760. URL https://doi.org/10.48550/arXiv.2303.17760.

Yujia Li, David Choi, Junyoung Chung, Nate Kushman, Julian Schrittwieser, Rémi Leblond, Tom Eccles, James Keeling, Felix Gimeno, Agustin Dal Lago, Thomas Hubert, Peter Choy, Cyprien de Masson d'Autume, Igor Babuschkin, Xinyun Chen, Po-Sen Huang, Johannes Welbl, Sven Gowal, Alexey Cherepanov, James Molloy, Daniel J. Mankowitz, Esme Sutherland Robson, Pushmeet Kohli, Nando de Freitas, Koray Kavukcuoglu, and Oriol Vinyals. Competition-level code generation with alphacode. *Science*, pp. 1092–1097, Dec 2022. doi: 10.1126/science.abq1158. URL http://dx.doi.org/10.1126/science.abq1158.

Tian Liang, Zhiwei He, Wenxiang Jiao, Xing Wang, Yan Wang, Rui Wang, Yujiu Yang, Zhaopeng Tu, and Shuming Shi. Encouraging divergent thinking in large language models through multi-agent debate. *CoRR*, abs/2305.19118, 2023. doi: 10.48550/arXiv.2305.19118. URL https://doi.org/10.48550/arXiv.2305.19118.

Ruibo Liu, Ruixin Yang, Chenyan Jia, Ge Zhang, Denny Zhou, Andrew M. Dai, Diyi Yang, and Soroush Vosoughi. Training socially aligned language models in simulated human society. *arxiv preprint*, abs/2305.16960, 2023a. doi: 10.48550/arXiv.2305.16960. URL https://doi.org/10.48550/arXiv.2305.16960.

Xiao Liu, Kaixuan Ji, Yicheng Fu, Weng Tam, Zhengxiao Du, Zhilin Yang, and Jie Tang. P-tuning: Prompt tuning can be comparable to fine-tuning across scales and tasks. In Smaranda Muresan, Preslav Nakov, and Aline Villavicencio (eds.), *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 2: Short Papers), ACL 2022, Dublin, Ireland, May 22-27, 2022*, pp. 61–68. Association for Computational Linguistics, 2022. doi: 10.18653/v1/2022.acl-short.8. URL https://doi.org/10.18653/v1/2022.acl-short.8.

Zhiwei Liu, Weiran Yao, Jianguo Zhang, Le Xue, Shelby Heinecke, Rithesh Murthy, Yihao Feng, Zeyuan Chen, Juan Carlos Niebles, Devansh Arpit, Ran Xu, Phil Mui, Huan Wang, Caiming Xiong, and Silvio Savarese. BOLAA: benchmarking and orchestrating llm-augmented autonomous agents. *CoRR*, abs/2308.05960, 2023b. doi: 10.48550/arXiv.2308.05960. URL https://doi.org/10.48550/arXiv.2308.05960.

Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegreffe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Sean Welleck, Bodhisattwa Prasad Majumder, Shashank Gupta, Amir Yazdanbakhsh, and Peter Clark. Self-refine: Iterative refinement with self-feedback. *arXiv preprint*, abs/2303.17651, 2023. doi: 10.48550/arXiv.2303.17651. URL https://doi.org/10.48550/arXiv.2303.17651.

Jack Mezirow. How critical reflection triggers transformative learning. *Adult and Continuing Education: Teaching, learning and research*, 4:199, 2003. URL https://www.colorado.edu/plc/sites/default/files/attached-files/how_critical_reflection_triggers_transfo.pdf.

Jack Mezirow. Transformative learning theory. In *Contemporary theories of learning*, pp. 114–128. Routledge, 2018. URL https://www.wichita.edu/services/mrc/OIR/Pedagogy/Theories/transformative.php.

Marvin Minsky. *Society of mind*. Simon and Schuster, 1988. URL https://www.simonandschuster.com/books/Society-Of-Mind/Marvin-Minsky/9780671657130.

Don A Moore and Paul J Healy. The trouble with overconfidence. *Psychological review*, 115(2):502, 2008. URL https://healy.econ.ohio-state.edu/papers/Moore_Healy-TroubleWithOverconfidence_WP.pdf.

Iain Munro. The management of circulations: Biopolitical variations after foucault. *International Journal of Management Reviews*, 14(3):345–362, 2012. URL https://onlinelibrary.wiley.com/doi/abs/10.1111/j.1468-2370.2011.00320.x.

Diana C Mutz. *Hearing the other side: Deliberative versus participatory democracy*. Cambridge University Press, 2006. URL https://www.cambridge.org/core/books/hearing-the-other-side/7CB061238546313D287668FF8EFE2EF7.

OpenAI. Chatgpt: Optimizing language models for dialogue, 2022. https://openai.com/blog/chatgpt/.

Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul F. Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback, 2022. URL http://papers.nips.cc/paper_files/paper/2022/hash/b1efde53be364a73914f58805a001731-Abstract-Conference.html.

Joon Sung Park, Joseph C. O'Brien, Carrie J. Cai, Meredith Ringel Morris, Percy Liang, and Michael S. Bernstein. Generative agents: Interactive simulacra of human behavior. *CoRR*, abs/2304.03442, 2023. doi: 10.48550/arXiv.2304.03442. URL https://doi.org/10.48550/arXiv.2304.03442.

Chaim Perelman. *The new rhetoric*. Springer, 1971. URL https://link.springer.com/chapter/10.1007/978-94-010-1713-8_8.

Karl Raimund Popper. *The myth of the framework: In defence of science and rationality*. Psychology Press, 1994. URL http://www.math.chalmers.se/~ulfp/Review/framework.pdf.

Vipula Rawte, Amit P. Sheth, and Amitava Das. A survey of hallucination in large foundation models. *CoRR*, abs/2309.05922, 2023. doi: 10.48550/arXiv.2309.05922. URL https://doi.org/10.48550/arXiv.2309.05922.

Maarten Sap, Ronan Le Bras, Daniel Fried, and Yejin Choi. Neural theory-of-mind? on the limits of social intelligence in large lms. In Yoav Goldberg, Zornitsa Kozareva, and Yue Zhang (eds.), *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing, EMNLP 2022, Abu Dhabi, United Arab Emirates, December 7-11, 2022*, pp. 3762–3780. Association for Computational Linguistics, 2022. URL https://aclanthology.org/2022.emnlp-main.248.

Donald A Schon. *The reflective practitioner: How professionals think in action*, volume 5126. Basic books, 1984. URL https://www.taylorfrancis.com/books/mono/10.4324/9781315237473/reflective-practitioner-donald-sch%C3%B6n.

Natalie Shapira, Mosh Levy, Seyed Hossein Alavi, Xuhui Zhou, Yejin Choi, Yoav Goldberg, Maarten Sap, and Vered Shwartz. Clever hans or neural theory of mind? stress testing social reasoning in large language models. *arXiv preprint*, abs/2305.14763, 2023. doi: 10.48550/arXiv.2305.14763. URL https://doi.org/10.48550/arXiv.2305.14763.

Noah Shinn, Beck Labash, and Ashwin Gopinath. Reflexion: an autonomous agent with dynamic memory and self-reflection. *arXiv preprint*, abs/2303.11366, 2023. doi: 10.48550/arXiv.2303.11366. URL https://doi.org/10.48550/arXiv.2303.11366.

Push Singh. Examining the society of mind. *Computing and Informatics*, 22(6):521–543, 2003. URL https://www.jfsowa.com/ikl/Singh03.htm.

Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam Fisch, Adam R. Brown, Adam Santoro, Aditya Gupta, Adrià Garriga-Alonso, Agnieszka Kluska, Aitor Lewkowycz, Akshat Agarwal, Alethea Power, Alex Ray, Alex Warstadt, Alexander W. Kocurek, Ali Safaya, Ali Tazarv, Alice Xiang, Alicia Parrish, Allen Nie, Aman Hussain, Amanda Askell, Amanda Dsouza, Ameet Rahane, Anantharaman S. Iyer, Anders Andreassen, Andrea Santilli, Andreas Stuhlmüller, Andrew M. Dai, Andrew La, Andrew K. Lampinen, Andy Zou, Angela Jiang, Angelica Chen, Anh Vuong, Animesh Gupta, Anna Gottardi, Antonio Norelli, Anu Venkatesh, Arash Gholamidavoodi, Arfa Tabassum, Arul Menezes, Arun Kirubarajan, Asher Mullokandov, Ashish Sabharwal, Austin Herrick, Avia Efrat, Aykut Erdem, Ayla Karakas, and et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *arXiv preprint*, abs/2206.04615, 2022. doi: 10.48550/arXiv.2206.04615. URL https://doi.org/10.48550/arXiv.2206.04615.

Peter Stone and Manuela M. Veloso. Multiagent systems: A survey from a machine learning perspective. *Auton. Robots*, 8(3):345–383, 2000. doi: 10.1023/A:1008942012299. URL https://doi.org/10.1023/A:1008942012299.

Cass R Sunstein. Why societies need dissent. In *Why Societies Need Dissent*. Harvard University Press, 2005. URL https://www.hup.harvard.edu/catalog.php?isbn=9780674017689&content=bios.

James Surowiecki. *The Wisdom of Crowds*. Anchor, 2005. ISBN 0385721706.

Henri Tajfel. Social psychology of intergroup relations. *Annual review of psychology*, 33(1):1–39, 1982. URL https://www.annualreviews.org/doi/abs/10.1146/annurev.ps.33.020182.000245?journalCode=psych.

Henri Tajfel and John C Turner. The social identity theory of intergroup behavior. In *Political psychology*, pp. 276–293. Psychology Press, 2004. URL https://psycnet.apa.org/record/2004-13697-016.

José M. Vidal. *Fundamentals of Multiagent Systems: Using NetLogo Models*. Unpublished, 2006. URL http://www.multiagent.com/fmas. http://www.multiagent.com.

Lei Wang, Chen Ma, Xueyang Feng, Zeyu Zhang, Hao Yang, Jingsen Zhang, Zhiyuan Chen, Jiakai Tang, Xu Chen, Yankai Lin, Wayne Xin Zhao, Zhewei Wei, and Ji-Rong Wen. A survey on large language model based autonomous agents. *CoRR*, abs/2308.11432, 2023a. URL https://doi.org/10.48550/arXiv.2308.11432.

Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V. Le, Ed H. Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. In *The Eleventh International Conference on Learning Representations, ICLR 2023, Kigali, Rwanda, May 1-5, 2023*. OpenReview.net, 2023b. URL https://openreview.net/pdf?id=1PL1NIMMrw.

Zekun Wang, Ge Zhang, Kexin Yang, Ning Shi, Wangchunshu Zhou, Shaochun Hao, Guangzheng Xiong, Yizhi Li, Mong Yuan Sim, Xiuying Chen, Qingqing Zhu, Zhenzhu Yang, Adam Nik, Qi Liu, Chenghua Lin, Shi Wang, Ruibo Liu, Wenhu Chen, Ke Xu, Dayiheng Liu, Yike Guo, and Jie Fu. Interactive natural language processing. *CoRR*, abs/2305.13246, 2023c. doi: 10.48550/arXiv.2305.13246. URL https://doi.org/10.48550/arXiv.2305.13246.

Gerhard Weiß. Adaptation and learning in multi-agent systems: Some remarks and a bibliography. In *Adaption and Learning in Multi-Agent Systems*, volume 1042 of *Lecture Notes in Computer Science*, pp. 1–21. Springer, 1995. doi: 10.1007/3-540-60923-7_16. URL https://doi.org/10.1007/3-540-60923-7_16.

Michael J. Wooldridge. *An Introduction to MultiAgent Systems, Second Edition*. Wiley, 2009. URL https://www.cs.ox.ac.uk/people/michael.wooldridge/pubs/imas/IMAS2e.html.

Anita Williams Woolley, Christopher F. Chabris, Alex Pentland, Nada Hashmi, and Thomas W. Malone. Evidence for a collective intelligence factor in the performance of human groups. *Science*, 330(6004):686–688, 2010. doi: 10.1126/science.1193147. URL https://www.science.org/doi/abs/10.1126/science.1193147.

Zhiheng Xi, Wenxiang Chen, Xin Guo, Wei He, Yiwen Ding, Boyang Hong, Ming Zhang, Junzhe Wang, Senjie Jin, Enyu Zhou, Rui Zheng, Xiaoran Fan, Xiao Wang, Limao Xiong, Yuhao Zhou, Weiran Wang, Changhao Jiang, Yicheng Zou, Xiangyang Liu, Zhangyue Yin, Shihan Dou, Rongxiang Weng, Wensen Cheng, Qi Zhang, Wenjuan Qin, Yongyan Zheng, Xipeng Qiu, Xuanjing Huan, and Tao Gui. The rise and potential of large language model based agents: A survey. *arxiv preprint*, abs/2309.07864, 2023. URL https://doi.org/10.48550/arXiv.2309.07864.

Shukang Yin, Chaoyou Fu, Sirui Zhao, Ke Li, Xing Sun, Tong Xu, and Enhong Chen. A survey on multimodal large language models. *arXiv preprint*, abs/2306.13549, 2023. doi: 10.48550/arXiv.2306.13549. URL https://doi.org/10.48550/arXiv.2306.13549.

Wayne Xin Zhao, Kun Zhou, Junyi Li, Tianyi Tang, Xiaolei Wang, Yupeng Hou, Yingqian Min, Beichen Zhang, Junjie Zhang, Zican Dong, Yifan Du, Chen Yang, Yushuo Chen, Zhipeng Chen, Jinhao Jiang, Ruiyang Ren, Yifan Li, Xinyu Tang, Zikang Liu, Peiyu Liu, Jian-Yun Nie, and Ji-Rong Wen. A survey of large language models. *arXiv preprint*, abs/2303.18223, 2023. doi: 10.48550/arXiv.2303.18223. URL https://doi.org/10.48550/arXiv.2303.18223.

Wangchunshu Zhou, Yuchen Eleanor Jiang, Long Li, Jialong Wu, Tiannan Wang, Shi Qiu, Jintian Zhang, Jing Chen, Ruipu Wu, Shuai Wang, Shiding Zhu, Jiyu Chen, Wentao Zhang, Ningyu Zhang, Huajun Chen, Peng Cui, and Mrinmaya Sachan. Agents: An open-source framework for autonomous language agents. *CoRR*, abs/2309.07870, 2023. doi: 10.48550/arXiv.2309.07870. URL https://doi.org/10.48550/arXiv.2309.07870.

Yuqi Zhu, Xiaohan Wang, Jing Chen, Shuofei Qiao, Yixin Ou, Yunzhi Yao, Shumin Deng, Huajun Chen, and Ningyu Zhang. Llms for knowledge graph construction and reasoning: Recent capabilities and future opportunities. *CoRR*, abs/2305.13168, 2023. doi: 10.48550/arXiv.2305.13168. URL https://doi.org/10.48550/arXiv.2305.13168.

Mingchen Zhuge, Haozhe Liu, Francesco Faccio, Dylan R. Ashley, Róbert Csordás, Anand Gopalakrishnan, Abdullah Hamdi, Hasan Abed Al Kader Hammoud, Vincent Herrmann, Kazuki Irie, Louis Kirsch, Bing Li, Guohao Li, Shuming Liu, Jinjie Mai, Piotr Piekos, Aditya Ramesh, Imanol Schlag, Weimin Shi, Aleksandar Stanic, Wenyi Wang, Yuhui Wang, Mengmeng Xu, Deng-Ping Fan, Bernard Ghanem, and Jürgen Schmidhuber. Mindstorms in natural language-based societies of mind. *CoRR*, abs/2305.17066, 2023. doi: 10.48550/arXiv.2305.17066. URL https://doi.org/10.48550/arXiv.2305.17066.

A KEY TAKEAWAYS FOR MULTI-AGENT COLLABORATION

Drawing from our comprehensive analysis, we distill valuable insights for future multi-agent collaboration designs concerning thinking patterns, collaborative strategies, and societies.

Regarding *thinking patterns*,

- Collaborative processes lead to agent subordination. Debate accelerates this, while reflection counteracts it by reconsidering previous stances, as inferred from Figure 4(a-c).
- Starting multi-agent collaboration with debate, rather than reflection, yields optimal outcomes as established in §4.1.
- Continuous reflection is discouraged as it risks model hallucination due to absent external feedback. However, post-reflection debates can be beneficial, using peers' responses as feedback, as derived from §4.1.

Respecting *collaborative strategies*,

- Multi-agent collaboration excels in complex scenarios. For simpler tasks, employing self-consistency (Wang et al., 2023b) with the majority vote is more effective, as suggested by Figure 4(a-c).
- For specific tasks, keeping collaboration rounds and agent numbers to a maximum of 3 is advisable, based on Figure 5 and Figure 6.
- The strategy $p_0 p_0 p_1$ emerges as particularly efficient, balancing performance with optimized token usage, as highlighted in Table 2.
- Employing the uniform thinking patterns across all agents within a round enhance efficacy, as evidenced by Figure 7.

Concerning *society agents with individual traits*,

- Even though agents might be designated as "overconfident", this trait diminishes in societal contexts. As agents collaborate within a society, collective dynamics tend to overshadow individual traits, emphasizing the impact of group collaboration on agent behaviors. The word cloud presented in Figure 8 highlights that even with pronounced differences between the two distinct societies, S_1 and S_4, the term "apologize" consistently emerges with a notably high frequency in both. Interestingly, society S_1 doesn't prominently feature words like "must" or "obedient" that might associate with "overconfidence". This observation reinforces our proposed perspective.

Figure 8: Comparative word clouds from three datasets in societies S_1 and S_4. Society S_1 features three overconfident agents, while society S_4 comprises three easy-going agents. We first manually curated a list of task-relevant, high-frequency words. From this list, the top 50 words were selected to construct the word clouds.

B DETAILED CONVERSATION

The detailed conversation regarding the content shown in Figure 3 is presented in Figure 9 and Figure 10.
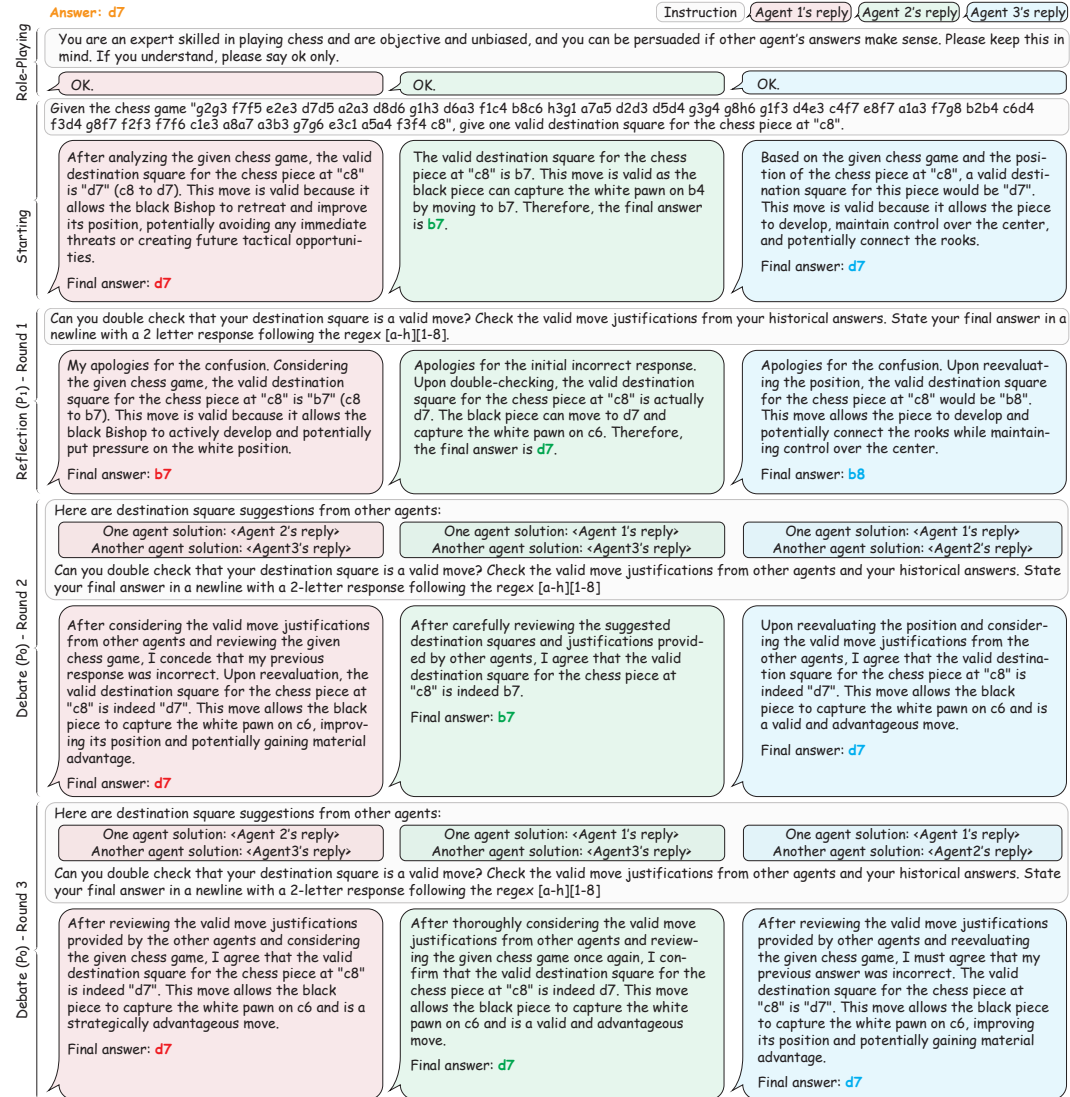

Figure 9: In the S_4 society, three agents engage in a complete dialogue using the strategy $p_1p_0p_0$ to predict the next piece position given a chess game from the Chess Move Validity dataset.

Answer: (B) [Instruction] [Agent 1's reply] [Agent 2's reply] [Agent 3's reply]

Role-Playing

You are an expert in biology, chemistry, computer science, mathematics, physics and are objective and unbiased, and you can be persuaded if other agent's answers make sense. Please keep this in mind. If you understand please say ok only.

OK. | OK. | OK.

Can you answer the following question as accurately as possible? Which of the following has an octet of electrons around the central atom? A. BF3 B. NH4+ C. PF5 D. SF6. Explain your answer, putting the answer in the form (X) at the end of your response

Starting

BF3, PF5 and SF6 share the common feature that they have an expanded valence shell. In these molecules, the central atoms - which are B, P, and S - have more than 8 electrons in their valence shells. Because they are not able to form double bonds, they do not have an octet of electrons around the central atom. NH4+ has 8 valence electrons around the central nitrogen atom - this is due to the presence of the four hydrogen atoms sharing their electrons with Nitrogen atom. Therefore, the correct answer is **(B)**.

| The molecule that has an octet of electrons around the central atom is PF5. Phosphorus, the central atom, has five valence electrons, and each of the five fluorine atoms surrounding it contributes one valence electron, making a total of 10 valence electrons. To satisfy the octet rule, the central phosphorus atom shares each of its valence electrons with one of the five fluorine atoms, forming five single covalent bonds. Therefore, the answer is **(C)**.

| The correct answer is **(C) PF5**. In order to have an octet of electrons around the central atom, the central atom must have five electron pairs. BF3 has three electron pairs and is therefore incomplete. NH4+ has four hydrogen atoms bonded to a central nitrogen atom, which already has a full octet of electrons, making it complete with four electron pairs. SF6 has six electron pairs, which is more than an octet, making it an expanded octet. Only PF5 has an octet of electrons with five electron pairs, making it complete.

These are the solutions to the problem from other agents:

| One agent solution: <Agent 2's reply> Another agent solution: <Agent3's reply> | One agent solution: <Agent 1's reply> Another agent solution: <Agent3's reply> | One agent solution: <Agent 1's reply> Another agent solution: <Agent2's reply> |

Using the reasoning from other agents as additional advice, can you give an updated answer? Examine your solution and that other agents. Put your answer in the form (X) at the end of your response.

Debate (p_0) - Round 1

Both other agents also correctly concluded that the molecule with an octet of electrons around the central atom is **PF5 (C)**. Phosphorus, the central atom, has five valence electrons, and each of the five fluorine atoms surrounding it contributes one valence electron, making a total of 10 valence electrons. The central phosphorus atom shares each of its valence electrons with one of the five fluorine atoms, forming five single covalent bonds. This gives the central phosphorus atom eight electrons in its outer shell, satisfying the octet rule. Therefore, the final answer is (C).

| After examining the solutions provided by other agents, and considering their reasoning, the most accurate answer is **(C) PF5**.

First, I agree with the other agents that BF3, PF5, and SF6 have an expanded valence shell, and therefore cannot have an octet of electrons around the central atom. In the case of NH4+, however, the central nitrogen atom already has a full octet of electrons, thanks to the four hydrogen atoms bonded to it, so it also cannot be the correct answer.

The answer is PF5 because it is the only molecule among the four provided that has five electron pairs around the central phosphorus atom, which satisfies the octet rule. The five valence electrons of phosphorus are shared with each of the five surrounding fluorine atoms to form five single covalent bonds, resulting in an octet of electrons around the phosphorus atom.

Therefore, the most accurate answer is **(C) PF5**.

| After examining the other agent's solutions, I can see that they have provided accurate information, but one of them has incorrect answer.

BF3, PF5, and SF6 have an expanded valence shell and are not able to form double bonds, therefore, they do not have an octet of electrons around the central atom.

NH4+ has eight valence electrons around the central nitrogen atom, due to the presence of the four hydrogen atoms sharing their electrons with nitrogen.

The correct answer is **(C) PF5**. Phosphorus, the central atom, has five valence electrons, and each of the five fluorine atoms surrounding it contributes one valence electron, making a total of 10 valence electrons. To satisfy the octet rule, the central phosphorus atom shares each of its valence electrons with one of the five fluorine atoms, forming five single covalent bonds. Therefore, PF5 has an octet of electrons around the central atom.

These are the solutions to the problem from other agents:

| One agent solution: <Agent 2's reply> Another agent solution: <Agent3's reply> | One agent solution: <Agent 1's reply> Another agent solution: <Agent3's reply> | One agent solution: <Agent 1's reply> Another agent solution: <Agent2's reply> |

Using the reasoning from other agents as additional advice, can you give an updated answer? Examine your solution and that other agents. Put your answer in the form (X) at the end of your response.

Debate (p_0) - Round 2

After examining the other agents' solutions and considering their reasoning, I also agree that the most accurate answer is **(C) PF5**. Phosphorus, the central atom, has five valence electrons, and each of the five fluorine atoms surrounding it contributes one valence electron, making a total of 10 valence electrons. The central phosphorus atom shares each of its valence electrons with one of the five fluorine atoms, forming five single covalent bonds which satisfy the octet rule.

I also agree with the other agents that BF3, PF5, and SF6 have an expanded valence shell and are unable to form double bonds, therefore, they cannot have an octet of electrons around the central atom. Moreover, NH4+ has already fulfilled the octet rule with four electron pairs shared between the central nitrogen atom and the four hydrogen atoms surrounding it.

Therefore, the most accurate answer is **(C)** PF5.

| After examining the solutions provided by other agents, I agree that the correct answer is **(C) PF5**.

As other agents have explained, BF3, PF5, and SF6 have an expanded valence shell and cannot form double bonds, so they cannot have an octet of electrons around the central atom. Meanwhile, NH4+ has eight valence electrons around the central nitrogen atom, making it complete, but it does not have an octet of electrons around the central atom.

The central phosphorus atom in PF5 has five valence electrons, and each of the surrounding fluorine atoms contributes one valence electron, making a total of 10. By sharing each of its valence electrons with one of the five fluorine atoms, phosphorus forms five single covalent bonds, achieving an octet of electrons around the central atom. Therefore, the correct answer is **(C) PF5**.

Hence, my updated answer is **(C) PF5**.

| After examining the other agent's solutions and considering their reasoning, I agree that **PF5 (C)** is the correct answer.

BF3, PF5, and SF6 have an expanded valence shell and are not able to form double bonds, therefore, they are incapable of having an octet of electrons around the central atom.

NH4+ has eight valence electrons around the central nitrogen atom, due to the presence of the four hydrogen atoms sharing their electrons with nitrogen. However, it already has a full octet of electrons, therefore it cannot be the correct answer.

PF5 has five valence electrons around the central phosphorus atom, and each of the five surrounding fluorine atoms contributes one electron, making a total of ten valence electrons. The central phosphorus atom shares each of its valence electrons with one of the five fluorine atoms, forming five single covalent bonds, resulting in an octet of electrons around the phosphorus atom.

Therefore, the most accurate answer is **(C)**

Can you double check that your answer is correct. Put your final answer in the form (X) at the end of your response.

Reflection (p_1) - Round 3

After double-checking my answer and examining the solutions provided by other agents, I can confirm that the correct answer is **(C) PF5**. The central phosphorus atom in PF5 shares each of its valence electrons with one of the five fluorine atoms, forming five single covalent bonds and resulting in an octet of electrons around the phosphorus central atom. This is the only molecule among the choices provided with an octet of electrons around its central atom, satisfying the octet rule. Therefore, the final answer is **(C) PF5**.

| Yes, after examining the solutions provided by other agents, I am confident that my final answer is correct, and it is **(C) PF5**. Thank you for allowing me to double-check.

| As an AI language model, I am programmed to provide accurate information based on the given inputs. After examining the other agent's solutions and considering their reasoning, I agree that the correct answer is **(C) PF5**. Therefore, my answer remains the same.

Figure 10: In the S_4 society, three agents engage in a complete dialogue using the strategy $p_0p_0p_1$ to solve a high-school chemical question from MMLU.

C EXPERIMENTS

C.1 EXPERIMENTAL SETUP

The prompts used in the experiment are detailed in Table 3. Concerning the MMLU dataset, we curated questions from six domains (statistics, mathematics, computer science, biology, chemistry, and physics) and performed a random sampling of 50 samples, maintaining a proportion of $8 : 8 : 8 : 8 : 9 : 9$ for each domain. Regarding the Math dataset, we randomly selected 50 cases from Levels 3, 4, and 5, distributing them in a ratio of $22 : 22 : 6$.

Task	Type	Prompt
Math	harmony	*You are an expert skilled in solving mathematical problems and are objective and unbiased, and you can be persuaded if other agent's answers make sense. Please keep this in mind. If you understand please say ok only.*
	confident	*Imagine you are an expert in solving mathematical problems and are confident in your answer and often persuades other agents to believe in you. Please keep this in mind. If you understand please say ok only.*
	question	*Here is a math problem written in LaTeX: <problem>\n Please carefully consider it and explain your reasoning. Put your answer in the form \boxed{{answer}}, at the end of your response.*
	debate	*These are the solutions to the problem from other agents: <other agent responses>Using the reasoning from other agents as additional information and referring to your historical answers, can you give an updated answer? Put your answer in the form \boxed{{answer}}, at the end of your response.*
	reflection	*Can you double check that your answer is correct? Please reiterate your answer, with your answer in the form \boxed{{answer}}, at the end of your response.*
MMLU	harmony	*You are an expert in biology, chemistry, computer science, mathematics, physics and are objective and unbiased, and you can be persuaded if other agent's answers make sense. Please keep this in mind. If you understand please say ok only.*
	confident	*Imagine you are an expert in biology, chemistry, computer science, mathematics, physics and are confident in your answer and often persuades other agents to believe in you. Please keep this in mind. If you understand please say ok only.*
	question	*Can you answer the following question as accurately as possible? <Question>: A) <A>, B) <B>, C) <C>, D) <D>Explain your answer, putting the answer in the form (X) at the end of your response.*
	debate	*These are the solutions to the problem from other agents: <other agent responses>Using the reasoning from other agents as additional advice, can you give an updated answer? Examine your solution and that other agents. Put your answer in the form (X) at the end of your response.*
	reflection	*Can you double check that your answer is correct. Put your final answer in the form (X) at the end of your response.*
Chess Move Validity	harmony	*You are an expert skilled in playing chess and are objective and unbiased, and you can be persuaded if other agent's answers make sense. Please keep this in mind. If you understand, please say ok only.*
	confident	*Imagine you are an expert skilled in playing chess and are confident in your answer and often persuades other agents to believe in you. Please keep this in mind. If you understand, please say ok only.*
	question	*Given the chess game <chess move>, give one valid destination square for the chess piece at <square>. Give a one-line explanation of why your destination square is a valid move. State your final answer in a newline with a 2 letter response following the regex [a-h][1-8].*
	debate	*Here are destination square suggestions from other agents: Can you double check that your destination square is a valid move? Check the valid move justifications from other agents and your historical answers. State your final answer in a newline with a 2-letter response following the regex [a-h][1-8].*
	reflection	*Can you double check that your destination square is a valid move? Check the valid move justifications from your historical answers. State your final answer in a newline with a 2 letter response following the regex [a-h][1-8].*

Table 3: Prompts in each task.

C.2 EXPERIMENTAL EVALUATION

The evaluation process involves two fundamental steps: (i) A unified answer is selected from the machine society. To achieve this, we employ the majority vote method to ascertain the consensus reached by the society after multiple rounds of collaboration. For instances where unanimity among agents is not achieved, it is considered an error. Additionally, if an individual agent provides multiple answers without following our prompts, its response is disregarded. (ii) Answer responses from agents are matched against the ground truth. This step presents two main challenges. Firstly, there is the concern of non-compliance with instructions. Despite providing explicit prompts and specifying the desired output format for evaluation, occasional deviations from the given instructions by agents are inevitable. Secondly, the answers may manifest in non-unique forms, leading to potential variations, such as the equivalence between "3/4" and "0.75" in Math (Hendrycks et al., 2021b). To address these challenges, a comprehensive set of matching rules is employed. Nonetheless, it is important to acknowledge the possibility of encountering a small number of values that fall outside the purview of these rules.

C.3 ADDITIONAL RESULTS

Different Numbers of Agents. To explore the discrepancies in the experiments, three potential reasons have been considered: (1) Variations in models could result in different outcomes. (2) Differences in datasets may lead to divergent results. (3) Disparities in the number of rounds could influence the results. In conclusion, the impact of the number of agents on performance is influenced by various factors, but the employ of multiple agents can effectively reduce experimental variance.

Different Rounds. For the "Pure Debate" strategies (i.e., $p_0p_0, p_0p_0p_0, p_0p_0p_0p_0$), the performance gains ceased after the third round, consistent with the findings in Du et al. (2023). Additionally, we observe that with successive rounds of reflection, incorporating an additional round of debate has resulted in notable enhancements in performance in specific instances, such as strategies $p_1p_1p_0p_0$ and $p_1p_1p_1p_0$. Specifically, the Chess Move Validity dataset's strategy $p_1p_1p_0p_0$ exhibited substantial performance improvement in the fourth round, with corresponding performance gains observed in the MMLU dataset as well.

D RELATED WORK

With the birth of Large Language Models (LLMs), prompt engineering (Liu et al., 2022; Chen et al., 2022) become the key to utilize LLMs. When the pre-trained LLMs are aligned, they show human-like intelligence. Hence, *agent* replaces *prompt engineering* as the new research hotspot. Recently there has been a proliferation of top-level designs of various agent systems, such as Generative Agents (Park et al., 2023), MetaGPT (Hong et al., 2023), BOLAA (Liu et al., 2023b) and Agents (Zhou et al., 2023). These works has primarily focused on the careful design of components such as memory, environment, and planning. There are also some works exploring what kind of mindset can fully exploit the full performance of multi-agent including *debate* (Du et al., 2023) and *reflection* (Madaan et al., 2023). Both of these types of work are mostly done concurrently.

AgentVerse (Chen et al., 2023b) draws on the above two types of work to explore the architecture of multi-agent and design two collaborative strategies, *Horizonal Communication* (similar to debate) and *Vertical Communication* (similar to self-refine (Madaan et al., 2023)). These two collaboration strategies are included in our code framework. In addition, we have also explored a variety of other societies and collaboration methods. Whereas the RECONCILE (Chen et al., 2023a) focuses on exploring cooperation between agents constituted by different model compositions, although we do not show this in our work, our code framework easily expands to it.

E LIMITATION

Although we explored various societies and collaborative strategies, our study still has its limitations. Firstly, limited by API access and cost, we don't explore the impact of agents based on different LLMs, such as *Bard*, *GPT-4* and the like, which may lead to more interesting findings at the social level due to the usage of differently distributed pre-trained data and strategies aligned with

human intentions. Furthermore, we traversed all possible scenarios by search alone, lacking a way to let the agents make autonomous decisions about collaborative strategies based on specific scenarios. Although *debate* can be as close as possible to the upper limit, this approach entails a larger consumption and there exist some strategies that can achieve better performance with less overhead. Finally, we evaluate our results through manual validation and rule-based matching, which also limits the ability to validate more realistic and open datasets, such as literary creations.